

Medical Insurance Cost Prediction

By Hakeem Lawrence

Dataset Link: <https://www.kaggle.com/datasets/mirichoi0218/insurance>

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

Data Inspection

```
In [2]: insurance = pd.read_csv('insurance.csv')
insurance.head(10)
```

```
Out[2]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
5	31	female	25.740	0	no	southeast	3756.62160
6	46	female	33.440	1	no	southeast	8240.58960
7	37	female	27.740	3	no	northwest	7281.50560
8	37	male	29.830	2	no	northeast	6406.41070
9	60	female	25.840	0	no	northwest	28923.13692

```
In [3]: insurance.shape
```

```
Out[3]: (1338, 7)
```

```
In [4]: insurance.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
```

```

4    smoker    1338 non-null    object
5    region    1338 non-null    object
6    charges    1338 non-null    float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB

```

```
In [5]: insurance.isna().sum()
```

```

Out[5]: age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64

```

Descriptive Statistics

```
In [6]: insurance.describe(include=[np.number]).round(3)
```

```

Out[6]:
```

	age	bmi	children	charges
count	1338.000	1338.000	1338.000	1338.000
mean	39.207	30.663	1.095	13270.422
std	14.050	6.098	1.205	12110.011
min	18.000	15.960	0.000	1121.874
25%	27.000	26.296	0.000	4740.287
50%	39.000	30.400	1.000	9382.033
75%	51.000	34.694	2.000	16639.913
max	64.000	53.130	5.000	63770.428

```
In [7]: insurance.columns
```

```

Out[7]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')

```

```

In [8]: #categorical column names
ins_cat_columns = ['sex', 'smoker', 'region']

```

```

In [9]: # count number of each category in columns

for i in ins_cat_columns:
    print(f'Count of value in {i}')
    print('-'*25)
    print((insurance[i].value_counts(normalize=True).round(3)))
    print(' '*15)

```

```

Count of value in sex
-----
male          0.505
female        0.495
Name: sex, dtype: float64

Count of value in smoker
-----
no            0.795
yes           0.205

```

Name: smoker, dtype: float64

Count of value in region

southeast 0.272

southwest 0.243

northwest 0.243

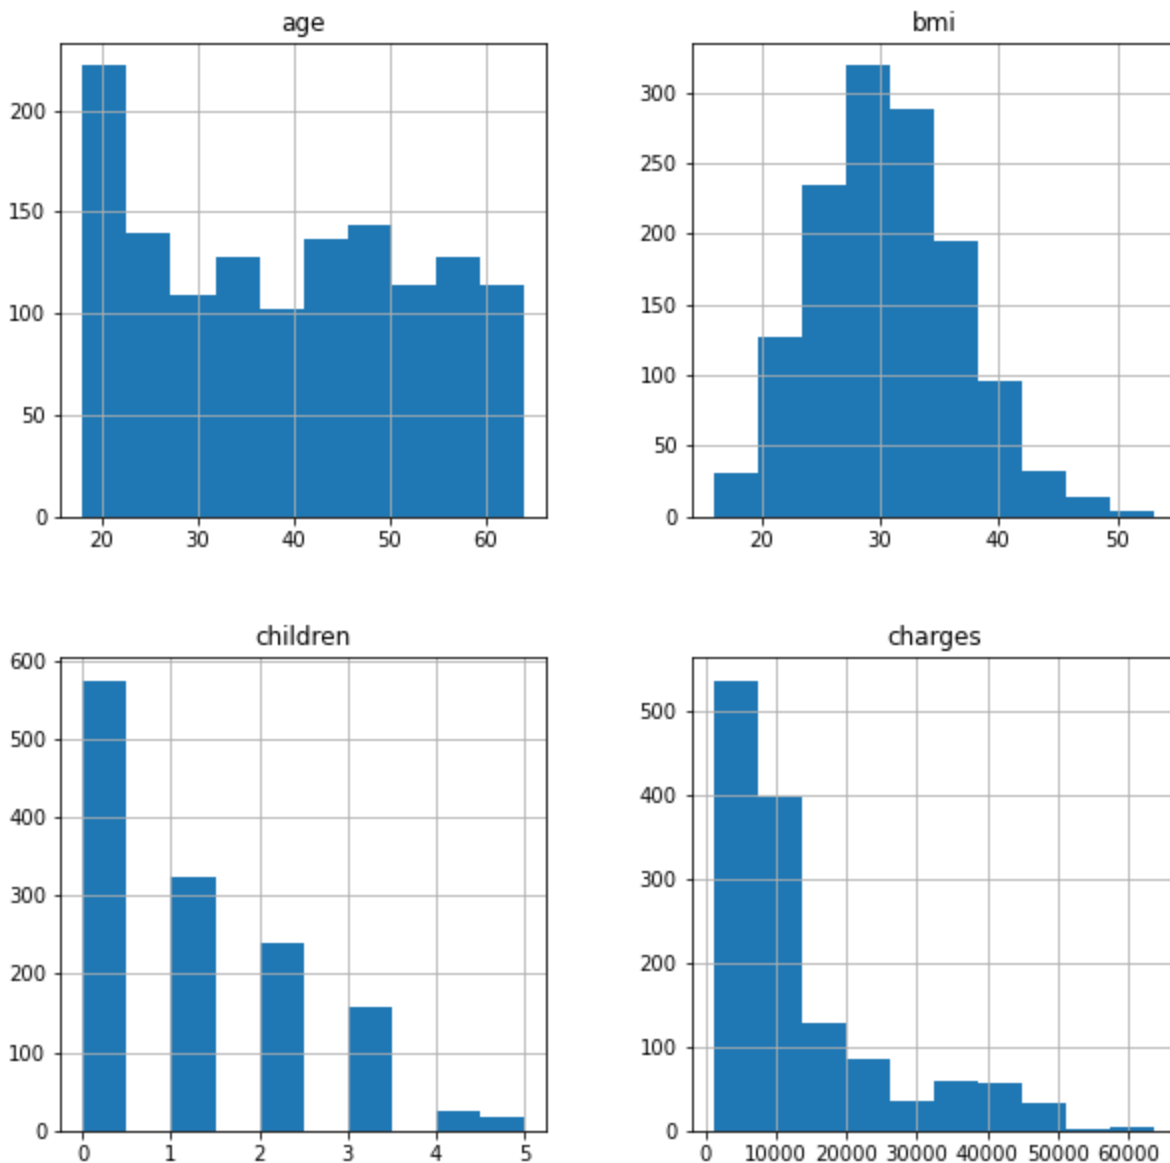
northeast 0.242

Name: region, dtype: float64

In [10]: *#distribution plots of numerical variables*

```
insurance.hist(figsize=(10,10))
```

Out[10]: array([[<AxesSubplot:title={'center':'age'}>,
<AxesSubplot:title={'center':'bmi'}>],
[<AxesSubplot:title={'center':'children'}>,
<AxesSubplot:title={'center':'charges'}>]], dtype=object)




In [11]: *#histogram of target variable*

```
sns.set(rc={'figure.figsize':(15,2)})  
sns.boxplot(data=insurance, x='charges',  
            showmeans=True,  
            meanprops={'marker':'o',  
                        'markerfacecolor':'white',  
                        'markeredgecolor':'black',  
                        'markersize': '10'})
```

```
plt.xlabel('charges', fontsize=16)
```

Out[11]: Text(0.5, 0, 'charges')






In [12]: *#creating function to display counts of categorical variables*

```
def counts(var):  
    plt.figure(figsize=(4,4))  
    for i in ins_cat_columns:  
        sns.countplot(data=insurance, x=var)  
        plt.title(f'Count by {var.title()}', size=15)
```

In [13]: *#showing counts of categorical variables*

```
for col in ins_cat_columns:  
    counts(col)
```





```
In [14]: #creating function to display distribution of variables by column

def distributions(variable):
    for i in ins_cat_columns:
        sns.displot(data=insurance, x=variable, hue=i,
                    height=4, aspect=2, kind='kde')


    plt.xlabel(variable, size=12)
    plt.ylabel('Count', size=12)
    plt.xticks(size=12)
    plt.title(f'Distribution of {variable.title()} by {i.title()}', size=16)
    plt.show()
```

```
In [15]: insurance.columns
```


```
Out[15]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```

```
In [16]: ins_num_columns= ['age', 'bmi', 'children', 'charges']
```


```
In [17]: for col in ins_num_columns:
          distributions(col)
```




Distribution of Age by Smoker




Distribution of Age by Region




Distribution of Bmi by Sex




Distribution of Bmi by Smoker




Distribution of Bmi by Region




Distribution of Children by Sex



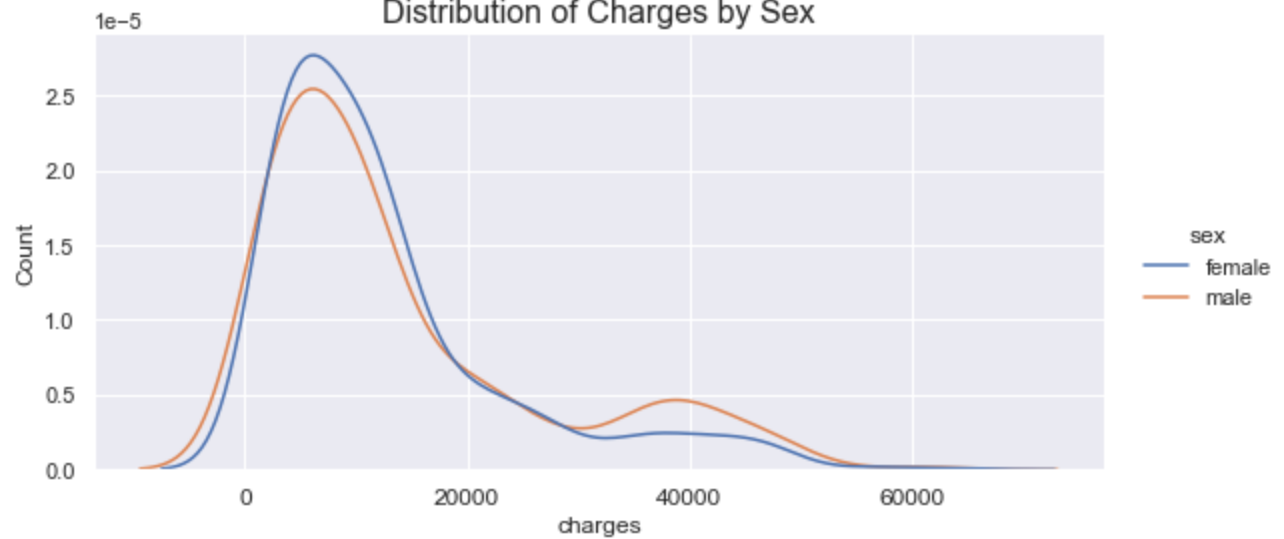
Distribution of Children by Smoker

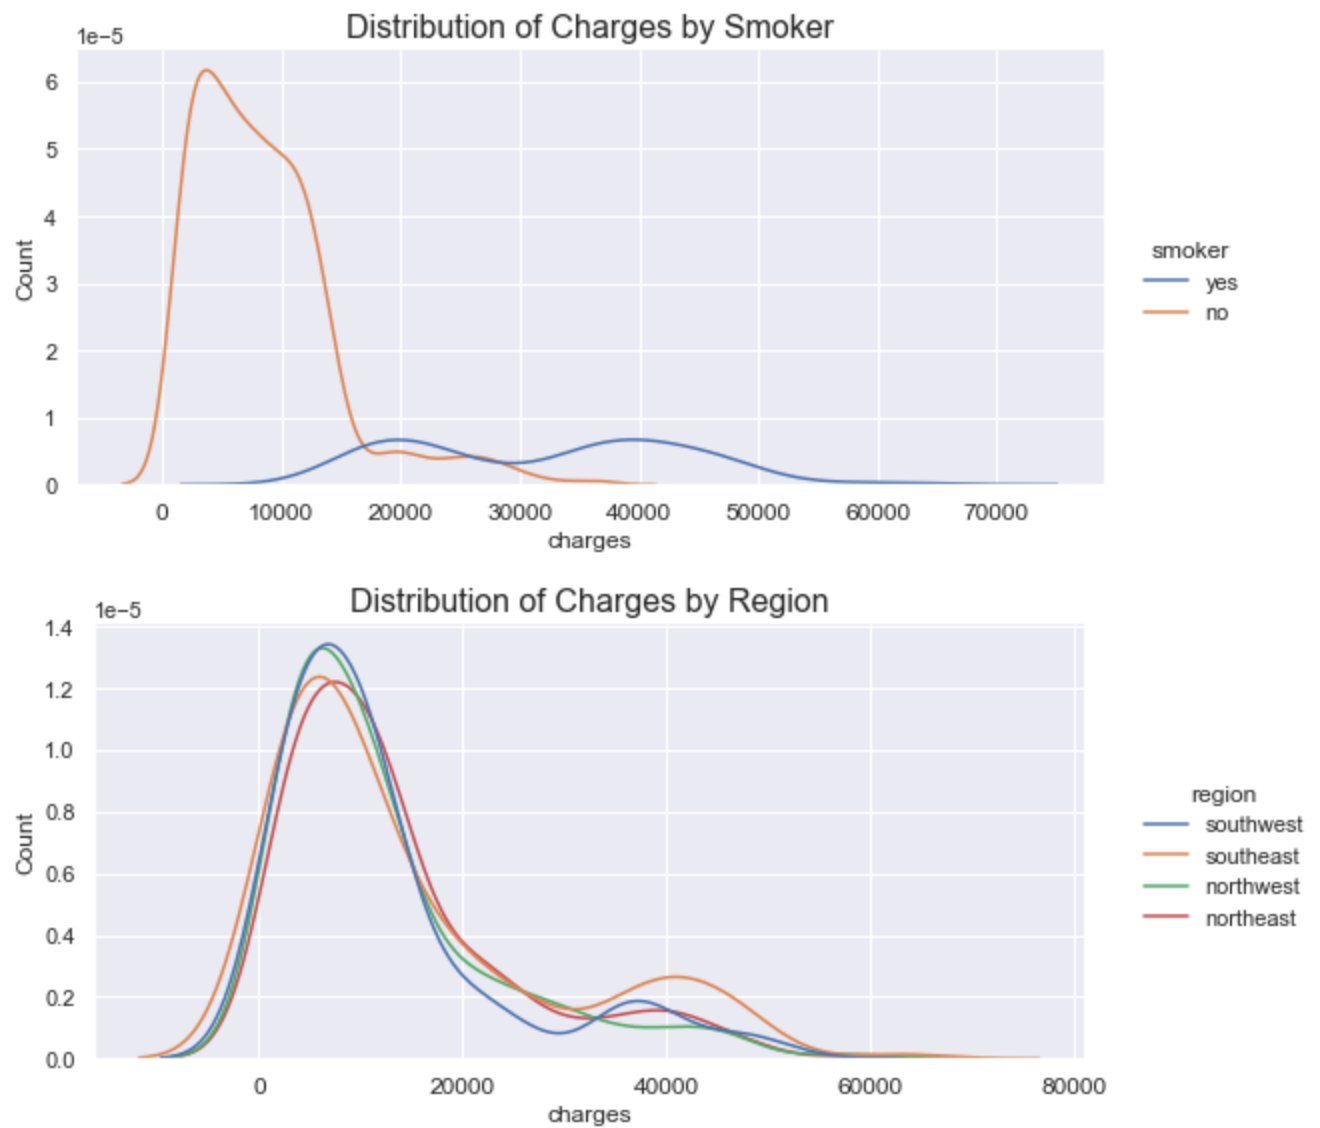


Distribution of Children by Region



Distribution of Charges by Sex






In [18]: *#creating function to display relationship between quant variables*

```
def scatter(var):
    if var != 'charges':
        plt.figure(figsize=(6,6))
        sns.scatterplot(x=var, y='charges', data=insurance)
        plt.title(f'Charges by {var.title()}', size=16)

for i in ins_num_columns:
    scatter(i)
```


Charges by Age



Charges by Bmi



Charges by Children




In [31]: *#viewing correlation between variables*

```
corr= insurance.corr()

plt.figure(figsize=(10,10))
sns.heatmap(corr,
            cbar=True,
            square=True, fmt='.3f',
            annot=True,
            annot_kws={'size':12},
            cmap='Blues')
```

Out[31]: <AxesSubplot:>



Pre-Processing

Encoding Categorical Features

```
In [20]: #encoding insurance column
insurance.replace({'sex':{'male':0,'female':1}}, inplace=True)

#encoding the smoker column
insurance.replace({'smoker':{'yes':0,'no':1}}, inplace=True)

#encoding region column
insurance.replace({'region':{'southwest':0,'southeast':1, 'northwest':2, 'northeast':3}})
```

Splitting Features from Target

```
In [21]: X= insurance.drop(columns='charges', axis=1)
Y= insurance['charges']
```

Training Model

```
In [22]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=.2, random_state=2)
```

```
In [23]: print(X.shape, X_train.shape, X_test.shape)
```

(1338, 6) (1070, 6) (268, 6)

```
In [24]: linearmodel = LinearRegression()
```

```
In [25]: linearmodel.fit(X_train, Y_train)
```

```
Out[25]: ▼ LinearRegression  
LinearRegression()
```

Model Evaluation

```
In [26]: #predicting on training data  
  
train_predict = linearmodel.predict(X_train)
```

```
In [27]: #R2 value  
  
r2_train = metrics.r2_score(Y_train, train_predict)  
print('R2 value:', r2_train)  
  
R2 value: 0.7519923667088932
```

```
In [28]: # predicting on test data  
  
test_predict = linearmodel.predict(X_test)
```

```
In [29]: #R2 value  
  
r2_test = metrics.r2_score(Y_test, test_predict)  
print('R2 value:', r2_test)  
  
R2 value: 0.7445422986536503
```

Prediction System

```
In [30]: input_data = (31,1,25.740,0,1,1) #charge = 3756.62160  
  
#changing input data to np array  
input_np = np.asarray(input_data)  
  
#reshape np array  
input_reshaped = input_np.reshape(1,-1)  
  
prediction = linearmodel.predict(input_reshaped).round(2)  
# print(prediction)  
  
charge = 3756.62160  
diff = (prediction - charge).round(2)  
# print(diff)  
  
print(f"The predicted insurance cost is ${prediction[0]}\n")  
print(f"The prediction is off by ${diff[0]} from the actual value")  
  
The predicted insurance cost is $3911.45  
  
The prediction is off by $154.83 from the actual value
```

```
C:\Users\hakee\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not
have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(
```